

ASSIGNMENT 7

September 15, 2021

1 ASSIGNMENT 7

2 SOUMYA GITE

3 SIRSS2276

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib as mpl
from sklearn.model_selection import train_test_split
```

```
[3]: data = pd.read_csv('Salary_Data.csv')
data.head(10)
```

```
[3]:   YearsExperience  Salary
0              1.1  39343.0
1              1.3  46205.0
2              1.5  37731.0
3              2.0  43525.0
4              2.2  39891.0
5              2.9  56642.0
6              3.0  60150.0
7              3.2  54445.0
8              3.2  64445.0
9              3.7  57189.0
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
```

```
dtypes: float64(2)
memory usage: 608.0 bytes
```

```
[5]: data.shape
```

```
[5]: (30, 2)
```

```
[6]: data.columns
```

```
[6]: Index(['YearsExperience', 'Salary'], dtype='object')
```

```
[7]: data.describe()
```

```
[7]:
```

	YearsExperience	Salary
count	30.000000	30.000000
mean	5.313333	76003.000000
std	2.837888	27414.429785
min	1.100000	37731.000000
25%	3.200000	56720.750000
50%	4.700000	65237.000000
75%	7.700000	100544.750000
max	10.500000	122391.000000

```
[8]: data.isnull().sum()
```

```
[8]: YearsExperience    0
Salary                0
dtype: int64
```

```
[9]: data.duplicated().sum()
```

```
[9]: 0
```

```
[10]: X = data.iloc[:, :-1].values
      Y = data.iloc[:, 1].values
```

```
[11]: class LinearRegression() :

    def __init__(self , learning_rate , iterations) :

        self.learning_rate = learning_rate

        self.iterations = iterations

    def fit(self , X , Y) :

        self.m, self.n = X.shape
```

```

self.W = np.zeros(self.n)

self.b = 0

self.X = X

self.Y = Y

# gradient descent learning

for i in range(self.iterations) :

    self.update_weights()

return self

def update_weights(self) :

    Y_pred = self.predict(self.X)

    # calculate gradients

    dW = -( 2 * (self.X.T).dot(self.Y - Y_pred )) / self.m

    db = - 2 * np.sum(self.Y - Y_pred ) / self.m

    # update weights

    self.W = self.W - self.learning_rate * dW

    self.b = self.b - self.learning_rate * db

    return self

def predict(self , X) :

    return X.dot(self.W) + self.b

```

```

[12]: X_train, X_test, Y_train, Y_test = train_test_split( X, Y, test_size = 1/3,
↳ random_state = 0 )

```

```

[13]: model = LinearRegression( iterations = 1000, learning_rate = 0.01 )
model.fit( X_train, Y_train )

```

```


[13]: <__main__.LinearRegression at 0x1d30ealf130>

```

```
[14]: Y_pred = model.predict( X_test )
print( "Predicted values ", np.round( Y_pred[:3], 2 ) )
print( "Real values      ", Y_test[:3] )
print( "Trained W         ", round( model.W[0], 2 ) )
print( "Trained b         ", round( model.b, 2 ) )
```

```
Predicted values [ 40594.69 123305.18  65031.88]
Real values      [ 37731. 122391.  57081.]
Trained W        9398.92
Trained b        26496.31
```

```
[15]: plt.figure(figsize=(8,6))
plt.scatter(X, Y)
plt.xlabel('YearsExperience')
plt.ylabel('Salary')
plt.show()
```



```
[16]: plt.scatter( X_test, Y_test, color = 'blue' )
plt.plot( X_test, Y_pred, color = 'orange' )
plt.title( 'Salary vs Experience' )
plt.xlabel( 'Years of Experience' )
plt.ylabel( 'Salary' )
```

```
plt.show()
```



```
[ ]:
```